

AEL - Product Requirements Document

Agentic Platform for the Handbook, Group 3: Tanishq Bhaiji, Nikola Hadzhiminov, Valentina Petrušić, Tanjin Rahman Rafi, Anisah Sakura Sutrisno

1. Problem and users

The primary problem addressed by the AEL platform is the challenge of turning raw field-research material and external sources into reliable, practical and evidence-based handbook chapters. Our team needs a repeatable process that allows non-coders to collect and assess sources, incorporate research, develop an AI-related use case, and produce a chapter that meets the handbook's quality requirements.

Our initial insight came from an employee at STMG, a large regional home-care provider. Our intended field-research partner, which is not yet formally confirmed, is Thuiszorg Totaal, a distinct home-care SME in Arnhem with approximately 70 staff. We treat STMG's account as a hypothesis to test at Thuiszorg Totaal, not as evidence about it. We expect this pattern may generalise beyond STMG's scale, though this is not yet independently confirmed and establishing whether it does is what our field research at Thuiszorg Totaal is for. The assumed problem is that clients are not reliably informed in advance of caretaker visits, which at STMG generates additional client calls. Whether the same pattern holds at Thuiszorg Totaal, and at what scale, is currently an interview-based observation, not quantified data. This is tested directly through our AIBS Internal Q1 ("How does Thuiszorg Totaal currently plan visits and communicate them to clients, and at which point does the information stop moving?") during the site visit.

Our research question is whether AI could help close this communication gap and, if so, how. The agentic platform supports this research by helping our team retrieve and evaluate relevant sources, analyse the information collected from Thuiszorg Totaal, and develop and document possible AI approaches. Its value compared with a human or a static wiki is that it can coordinate these research steps and repeatedly process and check information according to the defined workflow.

The platform will not automate decisions about caretaker schedules, client communication, or other operational decisions at Thuiszorg Totaal. It will also not make autonomous decisions about sensitive client information or replace human approval. The platform is intended to support our research and produce reliable handbook guidance, while any actual implementation of an AI solution at Thuiszorg Totaal would remain a separate decision for Thuiszorg Totaal.

Our five-person team is the primary user of the platform. The team has five distinct roles:

- Product Manager (Anisah Sakura Sutrisno): Writes the weekly research task brief with topic, priority, acceptance criteria that developers turn into CLI instructions, and gives final sign-off before publication.
- Developers (Tanishq Bhaiji, Nikola Hadzhiminov): Translate the PM's brief into CLI prompts for the resources intake and retrieval, run them, and refine the prompts when the Tester flags an issue.
- Tester (Valentina Petrušić): Checks whether each step of the platform produces the expected output, identifies errors or inconsistencies, and provides feedback through the agentic CLI before the workflow continues to the next step.
- Deployer (Tanjin Rahman Rafi, Nikola Hadzhiminov): Manage the final publication of the generated handbook chapters to the existing handbook site through the CL and verify that the output is accessible and correctly formatted.

The team works iteratively: each role reviews the relevant output before the next step is performed. The agentic CLI is therefore the main interface through which the team directs, tests and manages the platform.

Secondary users are Thuiszorg Totaal management and employees, who will read and potentially use the published handbook chapters, and the peer team, who will evaluate whether the platform works and whether the research in the chapters is well supported.

2. What it must let a user do

Raw material is uploadable to publish a chapter as an output. This will work as a sequence of steps that a non-coder can run one at a time and the team is checking the output before continuing:

1. The week's material, such as interview notes, CBS/KvK/sector sources, press articles and other relevant material, in Dutch or English can be drawn in.
2. Further sources in both Dutch and English for the week's research theme are retrieved.
3. Each source before it is used, considering factors such as publisher, recency, relevance and reliability is appraised.
4. Human review and edit: a team member checks the research, claims, language and recommendations before continuing.
5. Publish the approved chapter to the existing shared handbook.

The platform must allow the team to stop after each step, inspect the output, make corrections where necessary and only then continue to the next step. This ensures that the agent does not autonomously carry unverified information through the entire workflow.

3. Output qualities

A good handbook page must be clear, practical and easy to understand for Thuiszorg Totaal employees and management, including readers with little or no technical knowledge of AI.

To make these qualities measurable, each chapter will be reviewed against the following criteria:

- **Clear:** The chapter follows the fixed handbook structure, has a clearly stated problem and conclusion, and explains technical AI terminology in plain language when it is necessary. A reader should be able to identify the problem, proposed approach, requirements, risks and recommended next steps without needing additional technical knowledge.
- **Practical:** The chapter must provide concrete recommendations that are relevant to Thuiszorg Totaal . It must include five clearly defined first steps and identify the main requirements, resources, employee involvement and expected costs where reliable information is available.
- **Easy to understand:** The chapter should avoid unnecessary technical terminology and use short, direct explanations. The team will test this by having a person who was not involved in writing the chapter read it and explain the proposed AI use case and recommended first steps in their own words. If they cannot do this without substantial clarification, the chapter requires revision.
- **Evidence-based:** Claims about Thuiszorg Totaal must be traceable to our field research, while claims about AI, costs, technology or regulations must be supported by reliable external sources. Sources must be appraised before being used.
- **Responsible:** Each relevant chapter must address privacy, GDPR, risks, limitations and the role of employees. The chapter must clearly distinguish between what AI could assist with and what should remain under human responsibility.
- **Actionable:** The chapter must explain not only whether AI could help, but also what Thuiszorg Totaal would need to do to investigate or implement the proposed approach.

The final human review checks these criteria before publication. A chapter that does not meet the criteria is returned for revision rather than being published.

4. Known constraints

We direct the platform through an agentic CLI and do not hand-code the platform. The workflow must therefore be usable by our five-person team through language instructions and iterative review. The resulting chapters are published into the existing handbook.

Retrieval runs in both Dutch and English, because relevant Thuiszorg Totaal material, Chamber of Commerce information and regional coverage may be available primarily in Dutch.

Interview audio is transcribed in Dutch first and then translated into English. This allows the team to retain the original Dutch transcription so that translated claims and quotations can be checked against the original material. Team members who understand Dutch can manually verify or override a translation where necessary.

5. Out of scope, for now

The platform will not:

- automate the Thuiszorg Totaal interview itself;
- make decisions about caretaker schedules or client communication;
- autonomously use or interpret sensitive client information;
- generalise the research beyond the current SME case;
- automatically approve or publish a final chapter.

A human team member must always review and approve a chapter before it goes live.

6. Open questions and assumptions

- We assume that the "existing wiki and site" refers to the course's existing handbook. We still need to confirm how our team publishes to it (for example, through Git or an editor).
- Thuiszorg Totaal has not yet been formally confirmed as our field-research partner; the current PRD assumes that this partnership will be confirmed.
- The amount of interview access and client-data consent Thuiszorg Totaal provides remains open. This will determine how specific the illustrative AI use case can become.
- The evidence for Thuiszorg Totaal's reported communication problem currently comes from a specific guess. We still need to determine whether additional evidence, such as the number of client calls or the workload associated with them, is available.

Week 1 PRD - Response to the Socratic gate

The Socratic gate on our week 1 PRD raised four questions, which this revision addresses directly. Each of the five team members is now named against a specific role, and section 1 sets out how each role engages with the agentic CLI at a specific point in the workflow rather than describing the team's involvement in general terms. The claim that STMG's current planning system causes unnecessary workload is now explicitly presented as an interview-based hypothesis rather than an established fact, and is clearly separated from Thuiszorg Totaal, our actual research subject, whose experience this hypothesis is tested against through AIBS Internal Q1 during the site visit. Section 3 now defines six specific, testable criteria, clear, practical, easy to understand, evidence-based, responsible and actionable that determine whether a handbook page meets our quality bar, replacing the earlier general description with a standard we can apply consistently. The question of what an agentic platform specifically offers beyond a human researcher or a static wiki remains open at this stage. We expect to answer it concretely once our technical blueprint defines the platform's actual components, and we return to it in that document.