

AEL - Product Requirements Document

Agentic Platform for the Handbook, Group 3: Tanishq Bhaiji, Nikola Hadzhiminov, Valentina Petrušić, Tanjin Rahman Rafi, Anisah Sakura Sutrisno

Version 0.3

Date 24-09-2026

Reason for revision: Incorporates the week 2 Socratic questions and clarifies platform capabilities, source evaluation, privacy handling and measurable success criteria.

1. Problem and users

The primary problem addressed by the AEL platform is the challenge of turning raw field-research material and external sources into reliable, practical and evidence-based handbook chapters. Our team needs a repeatable process that allows non-coders to collect and assess sources, incorporate research, develop an AI-related use case, and produce a chapter that meets the handbook's quality requirements.

Our initial insight came from an employee at STMG, a large regional home-care provider. Thuiszorg Totaal is our provisional target SME and has not yet been formally confirmed as our field-research partner. The STMG account is treated as a hypothesis, not as evidence about Thuiszorg Totaal. The reported problem is that clients may not be reliably informed in advance about caretaker visits, potentially resulting in additional client calls. Whether this also occurs at Thuiszorg Totaal, and at what scale, will be investigated during the field research through AIBS Internal Q1.

Our research question is whether AI could help close this communication gap and, if so, how. The platform coordinates retrieval, appraisal, drafting and quality checking using predefined criteria and produces structured outputs for human review. It does not make final research or publication decisions.

The platform will not automate decisions about caretaker schedules, client communication, or other operational decisions at Thuiszorg Totaal. It will also not make autonomous decisions about sensitive client information or replace human approval. The platform is intended to support our research and produce reliable handbook guidance, while any actual implementation of an AI solution at Thuiszorg Totaal would remain a separate decision for Thuiszorg Totaal.

Our five-person team is the primary user of the platform. The team has five distinct roles:

- Product Manager (Anisah Sakura Sutrisno): Defines weekly research tasks, priorities and acceptance criteria and gives final sign-off.

- Developers (Tanishq Bhajji, Nikola Hadzhiminov): Translate requirements into CLI instructions, run the workflow and refine it when issues are identified.
- Tester (Valentina Petrušić): Checks workflow outputs against requirements and reports errors or inconsistencies.
- Deployer (Tanjin Rahman Rafi): Publishes approved chapters to the existing handbook site and verifies formatting/accessibility.

Secondary users are Thuiszorg Totaal management and employees, who read the handbook, and the peer team, who evaluate the platform and research.

2. What it must let a user do

The platform must turn raw research material into a handbook chapter through a sequence of human-reviewed steps that a non-coder can run one at a time.

1. Import the week's material, including interview notes/audio, CBS/KvK/sector sources and press material, in Dutch or English.
2. Retrieve additional relevant sources in both Dutch and English for the research theme.
3. Appraise each source before use using the defined source criteria. Calculate the source score, apply mandatory thresholds, and classify the source for use, rejection or human review.
4. Generate a chapter using the fixed handbook structure: SME-specific problem, plain-language explanation, CRISP-DM use case, tooling options, responsible-AI considerations, five first steps and sources.
5. Check the draft against the six chapter-quality criteria.
6. A team member reviews the research, claims, sources, language and recommendations and either approves the chapter or sends it back for revision.
7. Publish the approved chapter to the existing handbook site.

The team must be able to stop after every step, inspect and correct the output before continuing.

3. Output qualities

A good handbook page must be clear, practical and easy to understand for Thuiszorg Totaal employees and management, including readers with little AI knowledge. Each chapter is assessed against six criteria:

- **Clear:** follows the fixed structure and clearly identifies the problem, proposed approach, requirements, risks and next steps.
- **Practical:** contains relevant recommendations, five concrete first steps, requirements, employee involvement and costs where reliable information exists.
- **Easy to understand:** avoids unnecessary technical language and passes an external-reader comprehension test.

- **Evidence-based:** factual claims are traceable to field research or reliable external sources.
- **Responsible:** addresses privacy, GDPR, risks, limitations and human responsibility where relevant.
- **Actionable:** explains what the SME would need to do to investigate or implement the proposed approach.

Success metrics

Metric	Requirement
Quality pass rate	All 6 criteria pass before publication
Source traceability	100% of external factual claims have a traceable source
Understandability	Independent reader explains the problem, AI use case, risks and first steps without substantial clarification
Revision rate	Chapters returned for revision / chapters reviewed
Workflow intervention rate	Stages requiring correction / total stages run

Quality and traceability are checked for every chapter; revision and intervention rates are reviewed weekly.

Publication requires all six criteria to pass and human approval. Metrics are recorded per chapter and reviewed weekly.

4. Known constraints

The platform is directed through an agentic CLI without hand-coding and must be usable by the five-person team through natural-language instructions. Chapters are published to the existing handbook site; no new infrastructure is created.

Client-identifying information must not be sent to an external LLM API unless it has first passed the platform's defined sanitisation and validation process.

Retrieval runs in Dutch and English. Interview audio is transcribed in Dutch first and translated to English, while the original transcription is retained for verification.

Interview audio is transcribed in Dutch first and then translated into English. This allows the team to retain the original Dutch transcription so that translated claims and quotations can be checked against the original material. Team members who understand Dutch can manually verify or override a translation where necessary.

5. Out of scope, for now

Out of scope: automating the interview; making caretaker-scheduling or client-communication decisions; autonomously interpreting sensitive client

information; generalising beyond the current SME; and autonomous final approval/publication. Human approval is always required before publication.

6. Open questions and assumptions

- **Thuiszorg Totaal confirmation:** Provisional target SME; formal confirmation is still required.
- **ECD:** The ECD is unknown; this affects tooling and integration assumptions.
- **Interview/data access:** Access and client-data consent are unknown and determine the specificity of the use case.
- **Publishing:** The exact publication mechanism is still to be confirmed.
- **Problem evidence:** The initial problem observation comes from an STMG employee; further evidence such as call volume/workload data is still to be established.